

ECE-4300: Chapter-3: Arithmetic for Computers

FPU

Prof: Mohamed El-Hadedy

Email: mealy@cpp.edu

Office: 909-869-2594

Floating Point

Going beyond signed and unsigned integers, programming languages support numbers with fractions, which are called reals in mathematics.

Examples for reals:

1) pi: 3.14159265 (decimal)

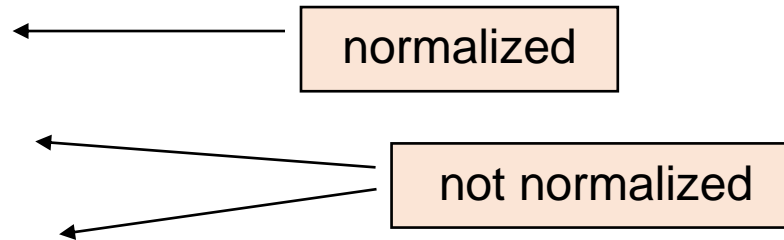
2) e: 2.71828

3) 0.000000001 $\rightarrow 1.0 \times 10^{-9}$

4) 3,155,760,000 or 3.15576×10^9

Floating Point

- Representation for non-integral numbers
 - Including very small and very large numbers
- Like scientific notation
 - -2.34×10^{56}
 - $+0.002 \times 10^{-4}$
 - $+987.02 \times 10^9$
- In binary
 - $\pm 1.xxxxxxx_2 \times 2^{yyyy}$
- Types float and double in C



Floating Point Standard

- Defined by IEEE Std 754-1985
- Developed in response to divergence of representations
 - Portability issues for scientific code
- Now almost universally adopted
- Two representations
 - Single precision (32-bit)
 - Double precision (64-bit)

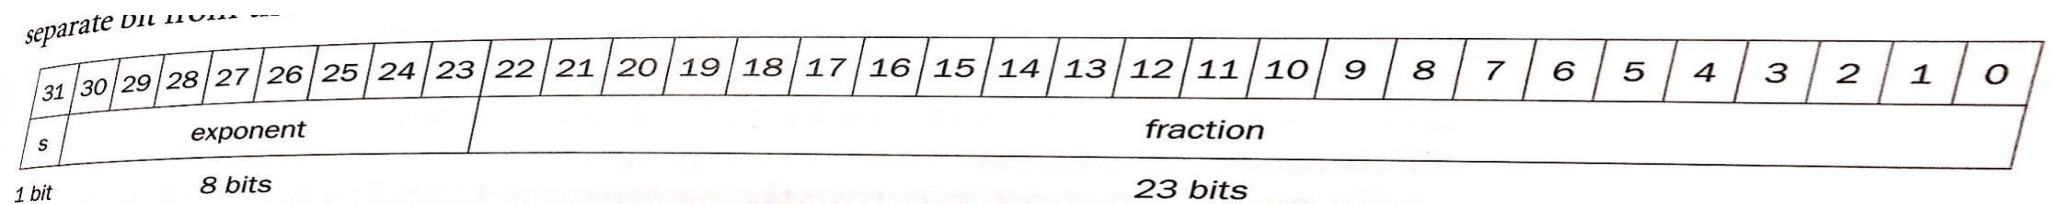
Floating-Point Representation

Precision	Base	Sign	Exponent	Significand
Single Precision	2	1	8	23+1
Double Precision	2	1	11	52+1

A finite number can also be represented by four integer components, a sign (s), a base (b), a significand (m), and an exponent (e). Then the numerical value of the number is evaluated as

$$(-1)^s \times m \times b^e \text{ ————— Where } m < |b|$$

Single-Precision



Floating Point Representation

- Numerical Form
 - $-1^s M 2^E$
 - Sign bit s determines whether number is negative or positive
 - Significand M normally a fractional value in range $[1.0, 2.0)$.
 - Exponent E weights value by power of two

- Encoding



- MSB is sign bit
- exp field encodes E
- frac field encodes M

Floating Point Precisions

- Encoding



- MSB is sign bit
 - exp field encodes E
 - frac field encodes M
- Sizes
 - Single precision: 8 exp bits, 23 frac bits
 - 32 bits total
 - Double precision: 11 exp bits, 52 frac bits
 - 64 bits total
 - Extended precision: 15 exp bits, 63 frac bits
 - Only found in Intel-compatible machines
 - Stored in 80 bits
 - 1 bit wasted

IEEE Floating-Point Format

single: 8 bits

double: 11 bits

single: 23 bits

double: 52 bits

S	Exponent	Fraction
---	----------	----------

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0 \Rightarrow non-negative, 1 \Rightarrow negative)
- Normalize significand: $1.0 \leq |\text{significand}| < 2.0$
 - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
 - Significand is Fraction with the “1.” restored
- Exponent: excess representation: actual exponent + Bias
 - Ensures exponent is unsigned
 - Single: Bias = 127; Double: Bias = 1023

Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
 - Exponent: 00000001
 \Rightarrow actual exponent = $1 - 127 = -126$
 - Fraction: 000...00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
 - exponent: 11111110
 \Rightarrow actual exponent = $254 - 127 = +127$
 - Fraction: 111...11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
 - Exponent: 000000000001
 \Rightarrow actual exponent = $1 - 1023 = -1022$
 - Fraction: 000...00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
 - Exponent: 111111111110
 \Rightarrow actual exponent = $2046 - 1023 = +1023$
 - Fraction: 111...11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

Floating-Point Precision

- Relative precision
 - all fraction bits are significant
 - Single: approx 2^{-23}
 - Equivalent to $23 \times \log_{10} 2 \approx 23 \times 0.3 \approx 6$ decimal digits of precision
 - Double: approx 2^{-52}
 - Equivalent to $52 \times \log_{10} 2 \approx 52 \times 0.3 \approx 16$ decimal digits of precision

Floating-Point Example

- Represent -0.75
 - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
 - $S = 1$
 - Fraction = $1000\dots00_2$
 - Exponent = $-1 + \text{Bias}$
 - Single: $-1 + 127 = 126 = 01111110_2$
 - Double: $-1 + 1023 = 1022 = 01111111110_2$
- Single: $10111111010000000000000000000000$
- Double: $10111111111010000000000000000000$

Floating-Point Example

- Represent -0.75
 - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
 - $S = 1$
 - Fraction = $1000\dots00_2$
 - Exponent = $-1 + \text{Bias}$
 - Single: $-1 + 127 = 126 = 01111110_2$
 - Double: $-1 + 1023 = 1022 = 01111111110_2$
- Single: $10111111010000000000000000000000_2$
- Double: $10111111111010000000000000000000_2$

Floating-Point Example (Converting Binary to Decimal Floating Point)

- What number is represented by the single-precision float

11000000101000...00

- $S = 1$
 - Fraction = $01000...00_2$
 - Exponent = $10000001_2 = 129$
- $x = (-1)^1 \times (1 + 01_2) \times 2^{(129 - 127)}$
 $= (-1) \times 1.25 \times 2^2$
 $= -5.0$

“Normalized” Numeric Values

- Condition
 - $\text{exp} \neq 000\dots 0$ and $\text{exp} \neq 111\dots 1$
- Exponent coded as *biased* value
$$E = \text{Exp} - \text{Bias}$$
 - *Exp* : unsigned value denoted by exp
 - *Bias* : Bias value
 - Single precision: 127 (*Exp*: 1...254, *E*: -126...127)
 - Double precision: 1023 (*Exp*: 1...2046, *E*: -1022...1023)
 - in general: $\text{Bias} = 2^{e-1} - 1$, where e is number of exponent bits
- Significand coded with implied leading 1
$$M = 1.\text{xxx}\dots\text{x}_2$$
 - xxx...x: bits of frac
 - Minimum when 000...0 ($M = 1.0$)
 - Maximum when 111...1 ($M = 2.0 - \square$)
 - Get extra leading bit for “free”

Normalized Encoding Example

- Value

Float $F = 15213.0$;

- $15213_{10} = 11101101101101_2 = 1.1101101101101_2 \times 2^{13}$

- Significand

$$M = \underline{1.1101101101101}_2$$

$$\text{frac} = \underline{1101101101101000000000}_2$$

- Exponent

$$E = 13$$

$$\text{Bias} = 127$$

$$\text{Exp} = 140 = 10001100_2$$

Floating Point Representation (Class 02):

Hex: 4 6 6 D B 4 0 0

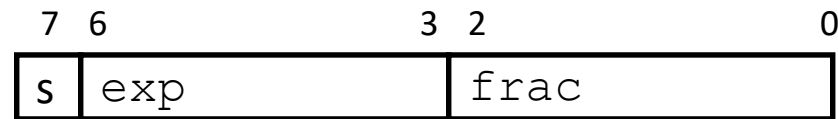
Binary: 0100 0110 0110 1101 1011 0100 0000 0000

140: 100 0110 0

15213: 1110 1101 1011 01

Tiny Floating Point Example

- 8-bit Floating Point Representation
 - the sign bit is in the most significant bit.
 - the next four bits are the exponent, with a bias of 7.
 - the last three bits are the frac
- Same General Form as IEEE Format
 - normalized, denormalized
 - representation of 0, NaN, infinity



Denormal Numbers


- Exponent = 000...0 \Rightarrow hidden bit is 0

$$x = (-1)^s \times (0 + \text{Fraction}) \times 2^{-\text{Bias}}$$

- Smaller than normal numbers
 - allow for gradual underflow, with diminishing precision
- Denormal with fraction = 000...0

$$x = (-1)^s \times (0 + 0) \times 2^{-\text{Bias}} = \pm 0.0$$

Two representations
of 0.0!



Floating-Point Addition

- Consider a 4-digit decimal example
 - $9.999 \times 10^1 + 1.610 \times 10^{-1}$
- 1. Align decimal points
 - Shift number with smaller exponent
 - $9.999 \times 10^1 + 0.016 \times 10^1$
- 2. Add significands
 - $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
- 3. Normalize result & check for over/underflow
 - 1.0015×10^2
- 4. Round and renormalize if necessary
 - 1.002×10^2

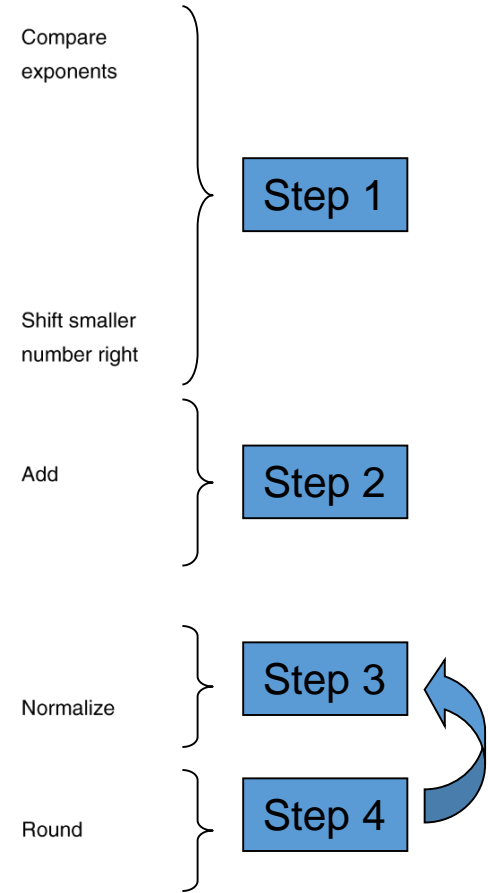
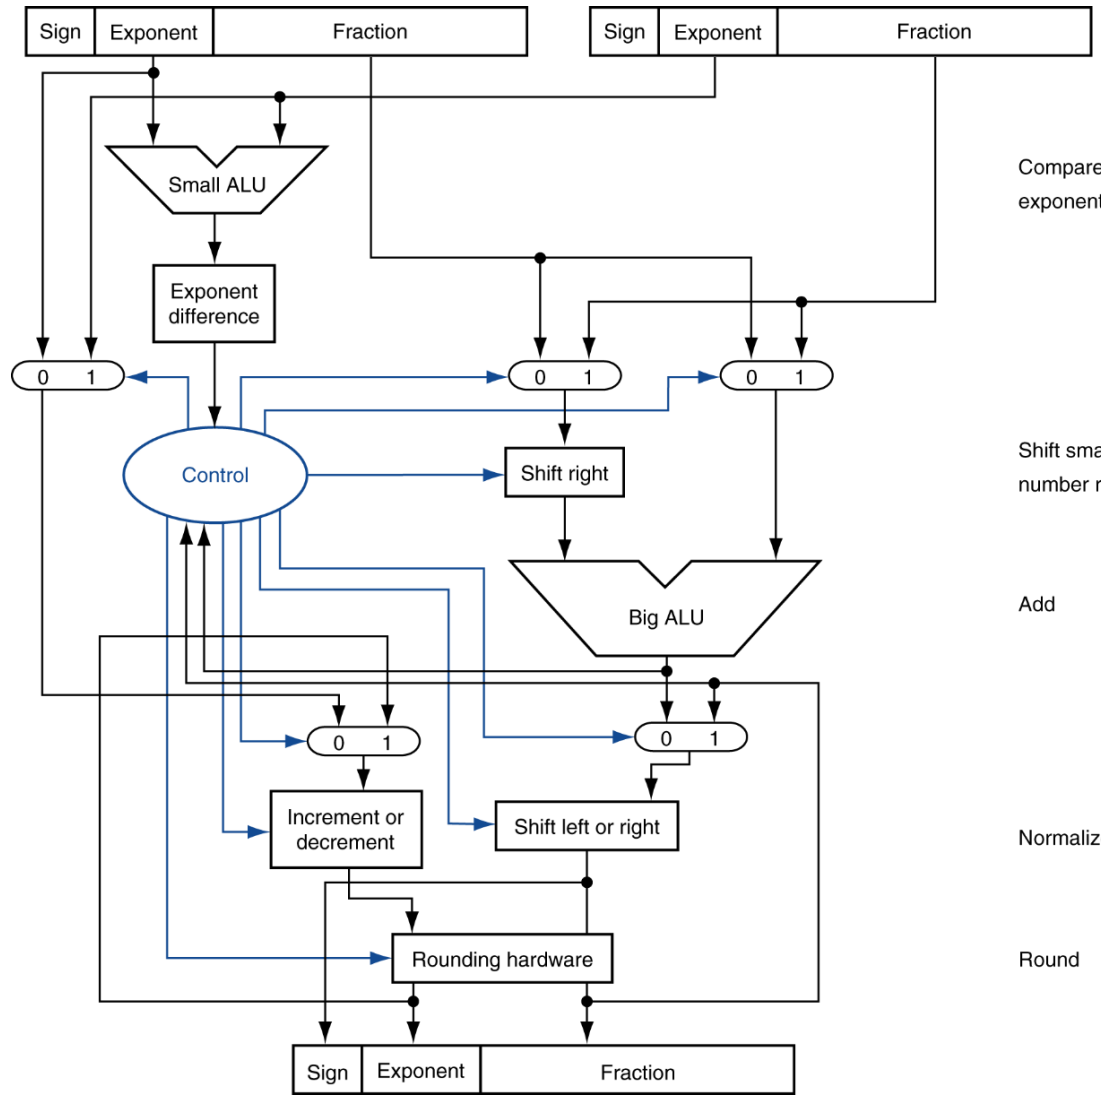
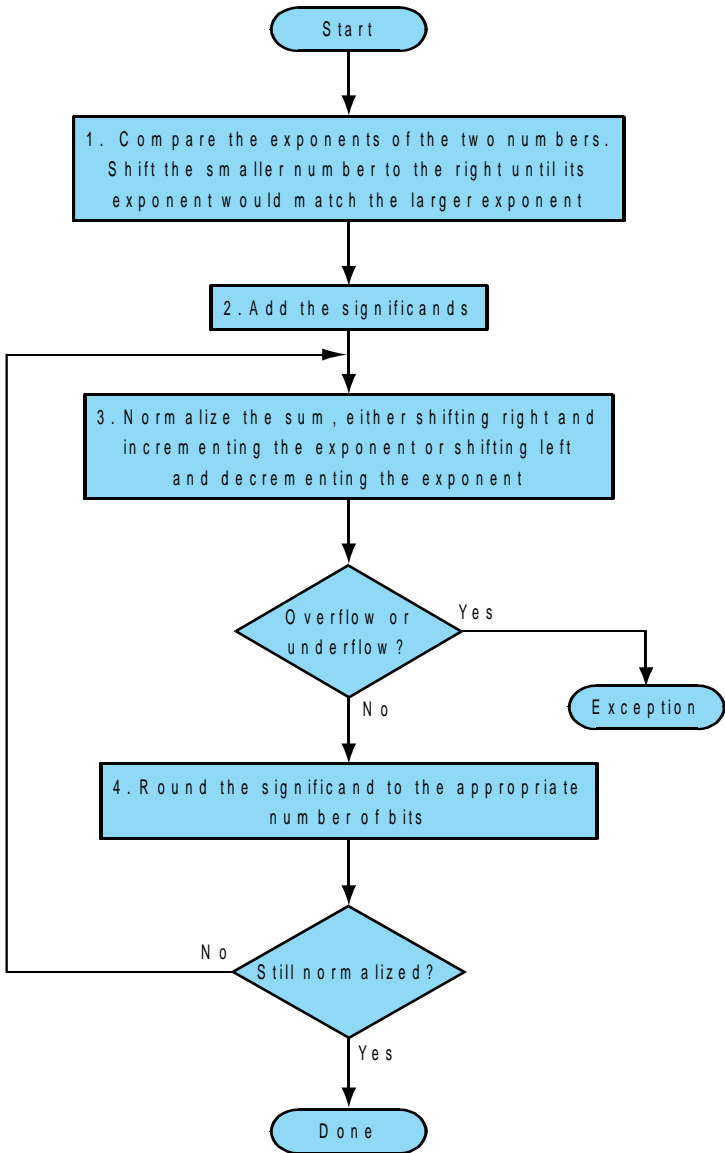
Floating-Point Addition

- Now consider a 4-digit binary example
 - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$ (0.5 + -0.4375)
- 1. Align binary points
 - Shift number with smaller exponent
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- 2. Add significands
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
 - $1.000_2 \times 2^{-4}$, with no over/underflow
- 4. Round and renormalize if necessary
 - $1.000_2 \times 2^{-4}$ (no change) = 0.0625

FP Adder Hardware

- Much more complex than integer adder
- Doing it in one clock cycle would take too long
 - Much longer than integer operations
 - Slower clock would penalize all instructions
- FP adder usually takes several cycles
 - Can be pipelined

FP Adder Hardware



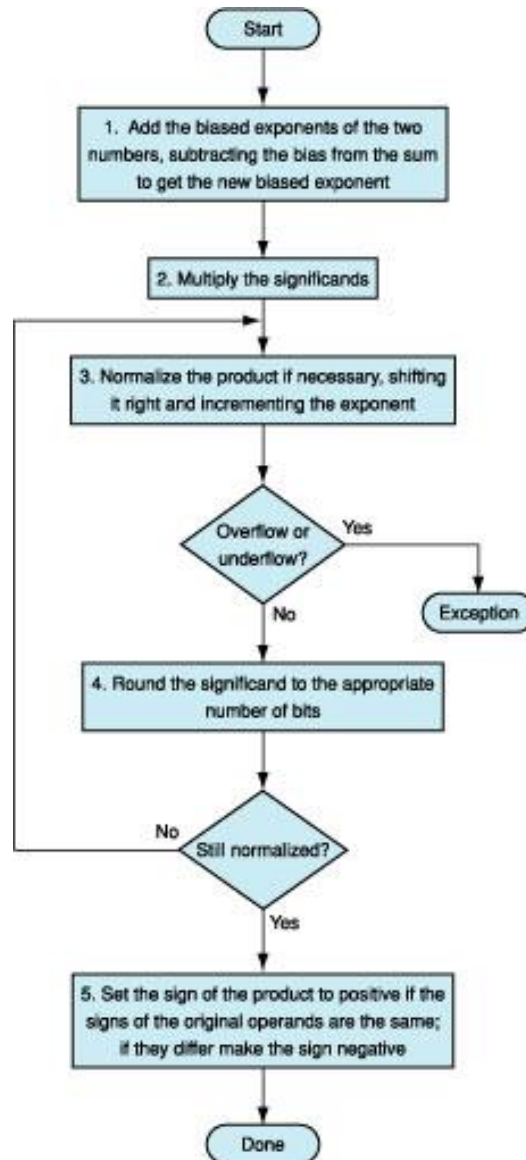
Floating-Point Multiplication

- Consider a 4-digit decimal example
 - $1.110 \times 10^{10} \times 9.200 \times 10^{-5}$
- 1. Add exponents
 - For biased exponents, subtract bias from sum
 - New exponent = $10 + -5 = 5$
- 2. Multiply significands
 - $1.110 \times 9.200 = 10.212 \Rightarrow 10.212 \times 10^5$
- 3. Normalize result & check for over/underflow
 - 1.0212×10^6
- 4. Round and renormalize if necessary
 - 1.021×10^6
- 5. Determine sign of result from signs of operands
 - $+1.021 \times 10^6$

Floating-Point Multiplication

- Now consider a 4-digit binary example
 - $1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2}$ (0.5×-0.4375)
- 1. Add exponents
 - Unbiased: $-1 + -2 = -3$
 - Biased: $(-1 + 127) + (-2 + 127) = -3 + 254 - 127 = -3 + 127$
- 2. Multiply significands
 - $1.000_2 \times 1.110_2 = 1.110_2 \Rightarrow 1.110_2 \times 2^{-3}$
- 3. Normalize result & check for over/underflow
 - $1.110_2 \times 2^{-3}$ (no change) with no over/underflow
- 4. Round and renormalize if necessary
 - $1.110_2 \times 2^{-3}$ (no change)
- 5. Determine sign: $+ve \times -ve \Rightarrow -ve$
 - $-1.110_2 \times 2^{-3} = -0.21875$

Floating-Point Multiplication



FP Arithmetic Hardware

- FP multiplier is of similar complexity to FP adder
 - But uses a multiplier for significands instead of an adder
- FP arithmetic hardware usually does
 - Addition, subtraction, multiplication, division, reciprocal, square-root
 - FP \leftrightarrow integer conversion
- Operations usually takes several cycles
 - Can be pipelined

Differences between FP and Integer Instructions

- # o Separate set of registers. (f0-f31)
- # o Separate instructions to operate on those registers ..
- # ... for arithmetic (add) ...
- # ... and loads and stores.
- # o Need for format conversion.

Differences between FP and Integer Instructions

Registers

#

MIPS Registers

#

MIPS has four sets of coprocessor registers.

The integer (GPR) registers are NOT one of the four sets.

Each set has 32 registers.

#

Co-processor 0: Processor and system control.

Co-processor 1: MIPS-32 floating-point

Co-processor 2: Reserved for special-purpose designs.

Co-processor 3: MIPS-64 floating-point

#

Differences between FP and Integer Instructions

Coprocessor 1 and The Floating-Point Registers

#

Floating point handled by co-processor 1, one of 4 co-processors.

#

MIPS floating point registers also called co-processor 1 registers.

MIPS floating point instructions called co-processor 1 instructions.

#

Registers named f0-f31.

Each register is 32 bits. (For MIPS-32)

Coprocessor 1 Instructions

Coprocessor 1 instructions use coprocessor 1 (FP) registers.

This includes instructions that do FP arithmetic ..

.. and other types of instructions.

Many coprocessor 1 instructions have "c1" in their names ..

.. for example, lwcl.

Coprocessor 1 **arithmetic** instructions **do not** have c1 in their names ..

.. but they include a /completer/ that indicates data type ..

.. for example, add.d, where ".d" is the completer.

Completers:

".s" Single-Precision Floating Point (32 bits)

".d" Double-Precision Floating Point (64 bits)

".w" Integer (32 bits)

Double-Precision Operands

#

In MIPS-32 (including MIPS-I) FP registers are 32 bits.

Instructions with ".d" operands get each operand from a pair of regs.

Register numbers must be even.

Example:

add.d \$f0, \$f2, \$f4 # { \$f0, \$f1 } = { \$f2, \$f3 } + { \$f4, \$f5 }

add.d \$f0, \$f2, \$f5 # ILLEGAL in MIPS 32, because f5 is odd.

Immediate Operands and Constant Registers

#

MIPS FP instructions do not take immediate values.

addi.s \$f0, \$f1, 2.3 # ILLEGAL, no immediate FP instructions.

There is no counterpart of integer register 0 (\$0 aka r0 aka \$zero).

add.s \$f0, \$f1, \$f2 # f0 has the sum,

Arithmetic Operations

#

Add double-precision (64-bit) operands.

#

MIPS: add.d \$f0, \$f2, \$f4

#

Load and Store

#

Load double (eight bytes into two consecutive registers).

#

MIPS: ldc1 \$f0, 8(\$t0)

#

#

Move Between Register Files (E.g., integer to FP)

#

MIPS: mtc1 \$t0, \$f0

Format Conversion

#

Convert from one format to another, e.g., integer to double.

#

MIPS: cvt.d.w \$f0, \$f2

Floating Point Condition Code Setting

```
#  
# Compare and set condition code.  
#  
# MIPS:      c.gt.d $f0, $f2  
#  
#
```

Conditional Branch

```
#  
# Branch on floating-point condition.  
#  
# MIPS: bclf TARGET # Branch coprocessor 1 [condition code] false.
```

FP Instructions in MIPS

- FP hardware is coprocessor 1
 - Adjunct processor that extends the ISA
- Separate FP registers
 - 32 single-precision: \$f0, \$f1, ... \$f31
 - Paired for double-precision: \$f0/\$f1, \$f2/\$f3, ...
 - Release 2 of MIPS ISA supports 32×64 -bit FP reg's
- FP instructions operate only on FP registers
 - Programs generally don't do integer ops on FP data, or vice versa
 - More registers with minimal code-size impact
- FP load and store instructions
 - lwcl, ldc1, swcl, sdcl
 - e.g., ldc1 \$f8, 32(\$sp)

FP Instructions in MIPS

FP Load and Store

MIPS

#

Load word in to coprocessor 1

lwc1 \$f0, 4(\$t4) # \$f0 = Mem[\$t4 + 4]

#

Load double in to coprocessor 1

ldc1 \$f0, 0(\$t4) # \$f0 = Mem[\$t4 + 0]; \$f1 = Mem[\$t4 + 4]

#

Store word from coprocessor 1.

swc1 \$f0, 4(\$t4) # \$f0 = Mem[\$t4 + 4]

#

Store double from coprocessor 1.

sdc1 \$f0, 0(\$t4) # Mem[\$t4 + 0] = \$f0; Mem[\$t4 + 4] = \$f1

FP Instructions in MIPS

Move Instructions

MIPS

#

Move to coprocessor 1

mtc1 \$t0, \$f0 # f0 = t0 Note that destination is *second* reg.

#

Move from coprocessor 1.

mfc1 \$t0, \$f0

FP Instructions in MIPS

#

Data Type Conversion

Convert between floating-point and integer formats.

NOTE: Values don't convert automatically, need to use these insns.

MIPS

#

To: s, d, w; From: s, d, w

#

cvt.TO.FROM fd, fs

#

cvt.d.w \$f0, \$f2 # \$f0 = convert_from_int_to_double(\$f2)

FP Instructions in MIPS

Setting Condition Codes

In preparation for a branch, set cond code based on FP comparison.

MIPS

#

Compare: fs COND ft

COND: eq, gt, lt, le, ge

FMT: s, d

#

c.COND.FMT fs, ft

Sets condition code to true or false.

Condition is false if either operand is not a number.

#

c.lt.d \$f0, \$f2 # CC = \$f0 < \$f2

bclt TARG # Branch if \$f0 < \$f2

nop

c.ge.d \$f0, \$f2 # CC = \$f0 >= \$f2

bclt TARG2 # Branch if \$f0 < \$f2

nop

Reachable?

FP Branches

MIPS

#

Branch insn specifies whether CC register true or false.

#

bclt TARG

nop

bclf TARG

nop

FP Instructions in MIPS

- Single-precision arithmetic
 - add.s, sub.s, mul.s, div.s
 - e.g., add.s \$f0, \$f1, \$f6
- Double-precision arithmetic
 - add.d, sub.d, mul.d, div.d
 - e.g., mul.d \$f4, \$f4, \$f6
- Single- and double-precision comparison
 - c.xx.s, c.xx.d (xx is eq, lt, le, ...)
 - Sets or clears FP condition-code bit
 - e.g. c.lt.s \$f3, \$f4
- Branch on FP condition code true or false
 - bclt, bclf
 - e.g., bclt TargetLabel

FP Example: °F to °C

- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

- fahr in \$f12, result in \$f0, literals in global memory space

- Compiled MIPS code:

```
f2c: lwcl $f16, const5($gp)  
     lwc2 $f18, const9($gp)  
     div.s $f16, $f16, $f18  
     lwcl $f18, const32($gp)  
     sub.s $f18, $f12, $f18  
     mul.s $f0, $f16, $f18  
     jr $ra
```

FP Example: Array Multiplication

- $X = X + Y \times Z$
 - All 32×32 matrices, 64-bit double-precision elements

- C code:

```
void mm (double x[],  
         double y[], double z[]) {  
    int i, j, k;  
    for (i = 0; i != 32; i = i + 1)  
        for (j = 0; j != 32; j = j + 1)  
            for (k = 0; k != 32; k = k + 1)  
                x[i][j] = x[i][j]  
                    + y[i][k] * z[k][j];  
}
```

- Addresses of x, y, z in \$a0, \$a1, \$a2, and
i, j, k in \$s0, \$s1, \$s2

FP Example: Array Multiplication

■ MIPS code:

	li	\$t1, 32	# \$t1 = 32 (row size/loop end)
	li	\$s0, 0	# i = 0; initialize 1st for loop
L1:	li	\$s1, 0	# j = 0; restart 2nd for loop
L2:	li	\$s2, 0	# k = 0; restart 3rd for loop
	sll	\$t2, \$s0, 5	# \$t2 = i * 32 (size of row of x)
	addu	\$t2, \$t2, \$s1	# \$t2 = i * size(row) + j
	sll	\$t2, \$t2, 3	# \$t2 = byte offset of [i][j]
	addu	\$t2, \$a0, \$t2	# \$t2 = byte address of x[i][j]
	l.d	\$f4, 0(\$t2)	# \$f4 = 8 bytes of x[i][j]
L3:	sll	\$t0, \$s2, 5	# \$t0 = k * 32 (size of row of z)
	addu	\$t0, \$t0, \$s1	# \$t0 = k * size(row) + j
	sll	\$t0, \$t0, 3	# \$t0 = byte offset of [k][j]
	addu	\$t0, \$a2, \$t0	# \$t0 = byte address of z[k][j]
	l.d	\$f16, 0(\$t0)	# \$f16 = 8 bytes of z[k][j]

...

FP Example: Array Multiplication

...

sll	\$t0, \$s0, 5	# \$t0 = i*32 (size of row of y)
addu	\$t0, \$t0, \$s2	# \$t0 = i*size(row) + k
sll	\$t0, \$t0, 3	# \$t0 = byte offset of [i][k]
addu	\$t0, \$a1, \$t0	# \$t0 = byte address of y[i][k]
l.d	\$f18, 0(\$t0)	# \$f18 = 8 bytes of y[i][k]
mul.d	\$f16, \$f18, \$f16	# \$f16 = y[i][k] * z[k][j]
add.d	\$f4, \$f4, \$f16	# f4=x[i][j] + y[i][k]*z[k][j]
addiu	\$s2, \$s2, 1	# \$k k + 1
bne	\$s2, \$t1, L3	# if (k != 32) go to L3
s.d	\$f4, 0(\$t2)	# x[i][j] = \$f4
addiu	\$s1, \$s1, 1	# \$j = j + 1
bne	\$s1, \$t1, L2	# if (j != 32) go to L2
addiu	\$s0, \$s0, 1	# \$i = i + 1
bne	\$s0, \$t1, L1	# if (i != 32) go to L1

Interpretation of Data

The BIG Picture

- Bits have no inherent meaning
 - Interpretation depends on the instructions applied
- Computer representations of numbers
 - Finite range and precision
 - Need to account for this in programs

Associativity

- Parallel programs may interleave operations in unexpected orders
 - Assumptions of associativity may fail

		$(x+y)+z$	$x+(y+z)$
x	-1.50E+38	0.00E+00	-1.50E+38
y	1.50E+38		1.50E+38
z	1.0	1.0	
		1.00E+00	0.00E+00

- Need to validate parallel programs under varying degrees of parallelism

x86 FP Architecture

- Originally based on 8087 FP coprocessor
 - 8×80 -bit extended-precision registers
 - Used as a push-down stack
 - Registers indexed from TOS: ST(0), ST(1), ...
- FP values are 32-bit or 64 in memory
 - Converted on load/store of memory operand
 - Integer operands can also be converted on load/store
- Very difficult to generate and optimize code
 - Result: poor FP performance

x86 FP Instructions

Data transfer	Arithmetic	Compare	Transcendental
FILD mem/ST(i) FISTP mem/ST(i) FLDPI FLD1 FLDZ	FIADDP mem/ST(i) FISUBRP mem/ST(i) FIMULP mem/ST(i) FIDIVRP mem/ST(i) FSQRT FABS FRNDINT	FICOMP FIUCOMP FSTSW AX/mem	FPATAN F2XMI FCOS FPTAN FPREM FPSIN FYL2X

- Optional variations
 - **I**: integer operand
 - **P**: pop operand from stack
 - **R**: reverse operand order
 - But not all combinations allowed

Streaming SIMD Extension 2 (SSE2)

- Adds 4×128 -bit registers
 - Extended to 8 registers in AMD64/EM64T
- Can be used for multiple FP operands
 - 2×64 -bit double precision
 - 4×32 -bit double precision
 - Instructions operate on them simultaneously
 - Single-Instruction Multiple-Data

Right Shift and Division

- Left shift by i places multiplies an integer by 2^i
- Right shift divides by 2^i ?
 - Only for unsigned integers
- For signed integers
 - Arithmetic right shift: replicate the sign bit
 - e.g., $-5 / 4$
 - $11111011_2 \gg 2 = 11111110_2 = -2$
 - Rounds toward $-\infty$
 - c.f. $11111011_2 \ggg 2 = 00111110_2 = +62$

Who Cares About FP Accuracy?

- Important for scientific code
 - But for everyday consumer use?
 - “My bank balance is out by 0.0002¢!” ☹
- The Intel Pentium FDIV bug
 - The market expects accuracy
 - See Colwell, *The Pentium Chronicles*

Concluding Remarks

- ISAs support arithmetic
 - Signed and unsigned integers
 - Floating-point approximation to reals
- Bounded range and precision
 - Operations can overflow and underflow
- MIPS ISA
 - Core instructions: 54 most frequently used
 - 100% of SPECINT, 97% of SPECFP
 - Other instructions: less frequent

Reference:

- **Book:** D. A. Patterson and J. L. Hennessy, **Computer Organization and Design: The Hardware/ Software Interface**, 5th Edition, San Mateo, CA: Morgan and Kaufmann. ISBN: 1-55860-604-1
- <https://www.mips.com/>
- <https://cs.stanford.edu/people/eroberts/courses/soco/projects/risc/mips/index.html>
- Professor El-Naga ECE-425 notes