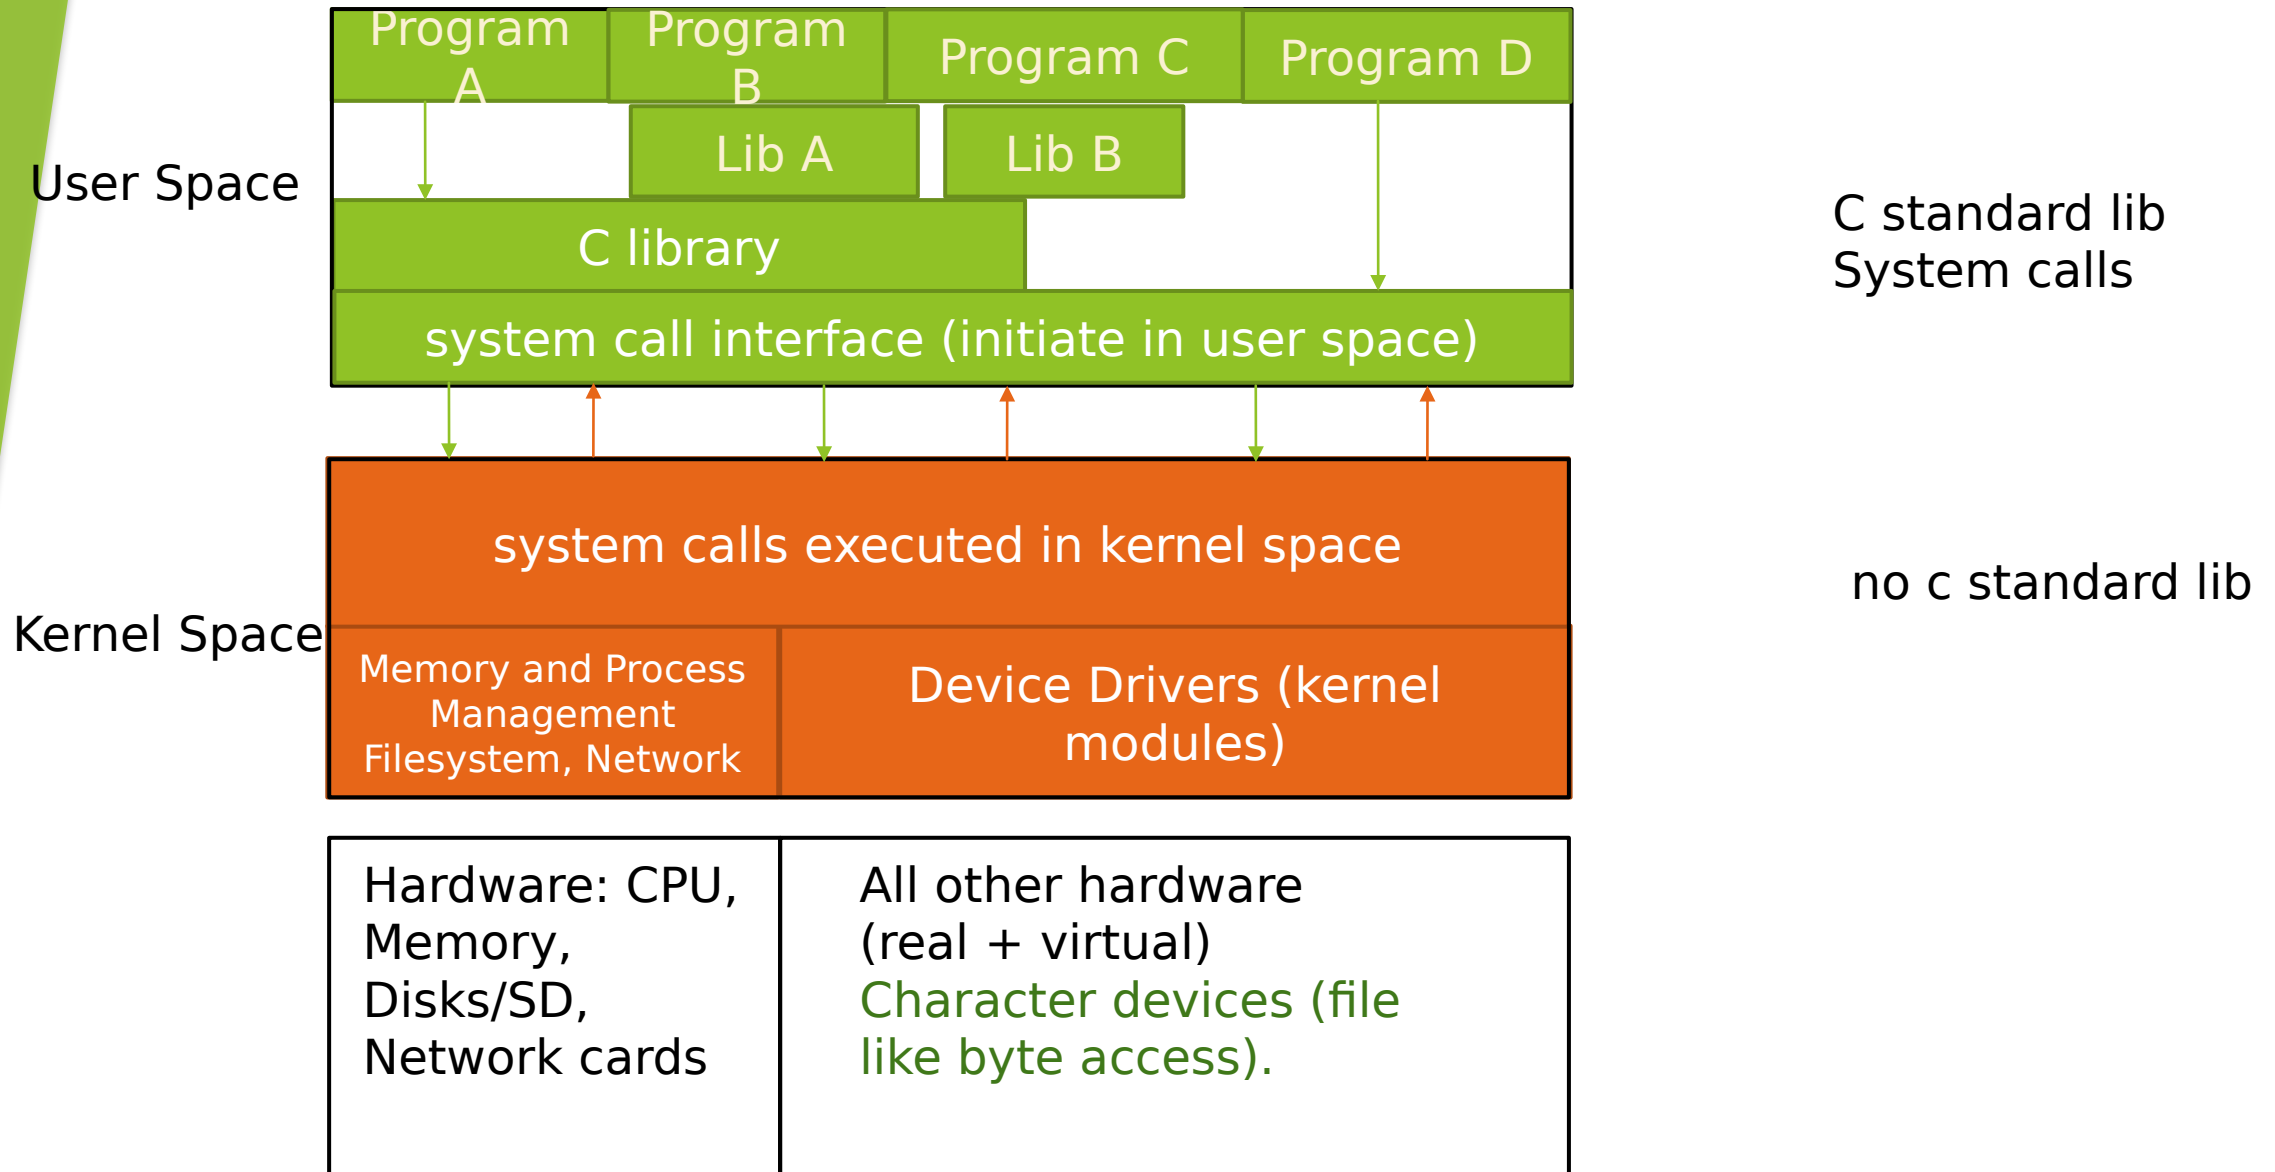


Operating Systems

Kernel Modules

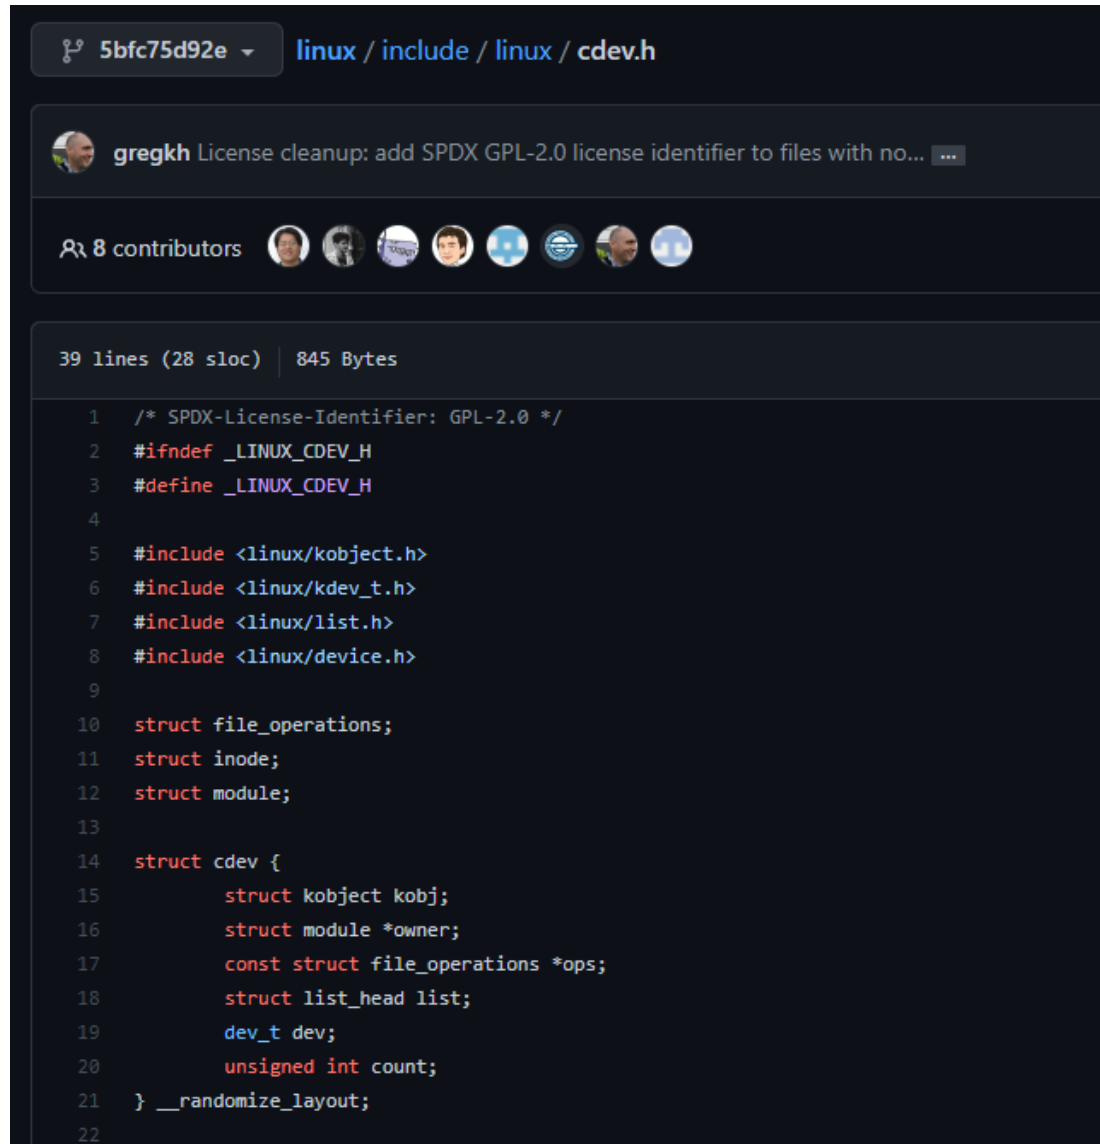
Kernel Space vs User Space



Kernel Modules

- Linux executable and linkable format (ELF)
 - Use gcc to compile/cross-compile
 - Runs in the kernel address space (high privilege)
 - Have use the extension “.ko”
- Loaded dynamically
 - Virtual
 - Control for some hardware device
- Manual load
 - insmod (will call the init callback)
 - Reserve resources, register callbacks
 - Create a “device file”/inode structure that includes a “file_operations” struct
 - rmmod (will call the exit callback)
 - Release resources (IRQ's, locks), free memory
- No main() entry point
 - init() and exit() points / callback functions
 - More callback functions can be added in the init() routine

Character Device



The screenshot shows a GitHub repository interface for the file `linux/include/linux/cdev.h`. At the top, the repository path is displayed as `linux / include / linux / cdev.h`. Below this, a commit by `gregkh` is shown with the message "License cleanup: add SPDX GPL-2.0 license identifier to files with no...". A row of 8 contributor avatars is visible. The file statistics indicate it is 39 lines (28 sloc) and 845 Bytes. The code content is as follows:

```
1  /* SPDX-License-Identifier: GPL-2.0 */
2  #ifndef _LINUX_CDEV_H
3  #define _LINUX_CDEV_H
4
5  #include <linux/kobject.h>
6  #include <linux/kdev_t.h>
7  #include <linux/list.h>
8  #include <linux/device.h>
9
10 struct file_operations;
11 struct inode;
12 struct module;
13
14 struct cdev {
15     struct kobject kobj;
16     struct module *owner;
17     const struct file_operations *ops;
18     struct list_head list;
19     dev_t dev;
20     unsigned int count;
21 } __randomize_layout;
22
```

Source: <https://github.com/torvalds/linux/blob/5bfc75d92efd494db37f5c4c173d3639d4772966/include/linux/cdev.h>

File Operations

```
struct file_operations {
    struct module *owner;
    loff_t(*llseek) (struct file *, loff_t, int);
    ssize_t(*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t(*aio_read) (struct kiocb *, char __user *, size_t, loff_t);
    ssize_t(*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t(*aio_write) (struct kiocb *, const char __user *, size_t,
                        loff_t);
    int (*readdir) (struct file *, void *, filldir_t);
    unsigned int (*poll) (struct file *, struct poll_table_struct *);
    int (*ioctl) (struct inode *, struct file *, unsigned int,
                 unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, struct dentry *, int datasync);
    int (*aio_fsync) (struct kiocb *, int datasync);
    int (*fasync) (int, struct file *, int);
    int (*lock) (struct file *, int, struct file_lock *);
    ssize_t(*readv) (struct file *, const struct iovec *, unsigned long,
                    loff_t *);
    ssize_t(*writev) (struct file *, const struct iovec *, unsigned long,
                     loff_t *);
    ssize_t(*sendfile) (struct file *, loff_t *, size_t, read_actor_t,
                       void __user *);
    ssize_t(*sendpage) (struct file *, struct page *, int, size_t,
                        loff_t *, int);
    unsigned long (*get_unmapped_area) (struct file *, unsigned long,
                                       unsigned long, unsigned long,
                                       unsigned long);
};
```

Basic Module

Find your version number using: `uname -a`
`ll /usr/src`

```
C test_004_pthreads.c  C test_006_pthreads_cond.c  M Makefile  C test_mod00.c 4 X
C test_mod00.c > MODULE_LICENSE
1  #include <linux/module.h>
2  #include <linux/kernel.h>
3  #include <linux/kern_levels.h>
4
5  MODULE_LICENSE("GPL");
6
7  int ece_init(void)
8  {
9      printk(KERN_INFO " ECE4310: Start here \n");
10     return 0;
11 }
12
13 void ece_end(void)
14 {
15     printk(KERN_INFO " ECE4310: End here \n");
16 }
17
18 module_init(ece_init);
19 module_exit(ece_end);
```

```
M Makefile  X  C test_mod00.c 4  C test_mod01.c  C read_write_mod.c
M Makefile
1  obj-m += test_mod00.o
2
3  KDIR=/usr/src/linux-headers-5.11.0-38-generic
4
5  CFLAGS = -D__KERNEL__ -DMODULE -I$(KDIR)/include -Wall
6
7  all:
8      $(MAKE) -C $(KDIR) M=$(shell pwd) $(KCONFIG) modules
9
10 clean:
11     $(MAKE) -C $(KDIR) M=$(shell pwd) clean
12     rm -f *.o
```

*Linux man-pages: full reference, not for learning. Check here when you know what you are looking for.
(<https://www.kernel.org/doc/man-pages/>)*

Basic Module

```
loniciuc@loniciuc-UVM:~$ ll /usr/src/
total 32
drwxr-xr-x  8 root root 4096 Nov 15 19:31 ./
drwxr-xr-x 14 root root 4096 Feb  9 2021 ../
drwxr-xr-x  7 root root 4096 Sep 29 18:53 linux-headers-5.11.0-37-generic/
drwxr-xr-x  7 root root 4096 Oct 20 06:21 linux-headers-5.11.0-38-generic/
drwxr-xr-x  7 root root 4096 Nov 15 19:31 linux-headers-5.11.0-40-generic/
drwxr-xr-x 24 root root 4096 Sep 29 18:53 linux-hwe-5.11-headers-5.11.0-37/
drwxr-xr-x 24 root root 4096 Oct 20 06:21 linux-hwe-5.11-headers-5.11.0-38/
drwxr-xr-x 24 root root 4096 Nov 15 19:30 linux-hwe-5.11-headers-5.11.0-40/
loniciuc@loniciuc-UVM:~$ uname -a
Linux loniciuc-UVM 5.11.0-38-generic #42~20.04.1-Ubuntu SMP Tue Sep 28 20:41:07 U
_64 x86_64 x86_64 GNU/Linux
```

- **make**
- *lsmod*
- *lsmod | grep test*
- **sudo insmod test_mod00.ko**
- *lsmod | grep test*
- *dmesg | tail*
- **sudo rmmod test_mod00**
- *dmesg | tail*

Globals

```
7  MODULE_LICENSE("GPL");
8
9  #define ECE_BUF_SIZE 4096
10
11  static char ece_buffer[ECE_BUF_SIZE];
12  int isReg;
13  int major;
14  int ece_offset_w;
15  int ece_offset_r;
16  int ece_size;
```


Write

```
43 static ssize_t ece_write(struct file *fp, const char *buf, size_t count, loff_t *op)
44 {
45     int ret = 0;
46     if(ece_offset_w + count >= ECE_BUF_SIZE)
47     {
48         printk(KERN_INFO " ECE4310: Write OverFlow. Abort. \n");
49         return -1;
50     }
51
52     printk(KERN_INFO " ECE4310: Copy from user. \n");
53     ret = copy_from_user(&ece_buffer[ece_offset_w], buf, count);
54     if(ret != 0)
55     {
56         printk(KERN_INFO " ECE4310: ERR copy from user. \n");
57         return -1;
58     }
59     ece_offset_w = ece_offset_w + count;
60
61     return count;
62 }
```

Read

```
19 static ssize_t ece_read(struct file *fp,
20     char *buf,          /* to fill with data for user */
21     size_t count,       /* how much data to send      */
22     loff_t *offset)     /* take the data from offset */
23 {
24     int ret;
25     if(ece_offset_r + count >= ECE_BUF_SIZE)
26     {
27         printk(KERN_INFO " ECE4310: Read OverFlow. Abort. \n");
28         return -1;
29     }
30     printk(KERN_INFO " ECE4310: Copy to user. \n");
31     ret = copy_to_user(buf, &ece_buffer[ece_offset_r], count);
32     if(ret != 0)
33     {
34         printk(KERN_INFO " ECE4310: ERR copy to user. \n");
35         return -1;
36     }
37
38     ece_offset_r = ece_offset_r + count;
39
40     return count;
41 }
```

Updated init

```
64 static struct file_operations ece_fops =
65 {
66     .read = ece_read,
67     .write = ece_write,
68 };
69
70 int ece_init(void)
71 {
72     int ret = 0;
73     major = register_chrdev(0, "test_mod01", &ece_fops);
74     ece_offset_w = 0;
75     ece_offset_r = 0;
76     ece_size = 0;
77
78     if(major < 0)
79     {
80         isReg = 0;
81         printk(KERN_INFO " ECE4310: Start FAIL \n");
82     }
83     else
84     {
85         isReg = 1;
86         printk(KERN_INFO " ECE4310: Start here \n");
87     }
88     return 0;
89 }
```

Updated end function

```
91 void ece_end(void)
92 {
93     if(isReg) //if(major >= 0)
94     {
95         unregister_chrdev(major, "test_mod01");
96     }
97     printk(KERN_INFO " ECE4310: End here \n");
98 }
99
100 module_init(ece_init);
101 module_exit(ece_end);
```

Kernel Modules

- Major/Minor Numbers

```
loniciuc@loniciuc-UVM:c$ cat /proc/devices | grep test
237 test_mod01
```

```
loniciuc@loniciuc-UVM:c$ sudo mknod /dev/test_mod01 c 237 1
```

```
loniciuc@loniciuc-UVM:c$ ll /dev/ | grep test
crw-r--r--  1 root    root    237,   1 Nov 14 19:18 test_mod01
```

```
loniciuc@loniciuc-UVM:c$ sudo chmod 0777 /dev/test_mod01
loniciuc@loniciuc-UVM:c$ ll /dev/test_mod01
crwxrwxrwx 1 root root 237, 1 Nov 14 19:18 /dev/test_mod01
```

Kernel Modules

```
loniciuc@loniciuc-UVM:c$ echo "HI CLASS.... is it late?" > /dev/test_mod01
loniciuc@loniciuc-UVM:c$ echo "maybe...." > /dev/test_mod01
loniciuc@loniciuc-UVM:c$ echo "SOME LOG MESSAGE IN HERE" > /dev/test_mod01
loniciuc@loniciuc-UVM:c$ ./a.out
```

```
result: First Tests
```

```
loniciuc@loniciuc-UVM:c$ ./a.out
```

```
result: New message
```

```
loniciuc@loniciuc-UVM:c$ ./a.out
```

```
result: HI CLASS..
```

```
loniciuc@loniciuc-UVM:c$ ./a.out
```

```
result: .. is it l3
```

```
loniciuc@loniciuc-UVM:c$ ./a.out
```

```
result: ate?
```

```
maybeR
```

```
loniciuc@loniciuc-UVM:c$
```

Kernel Modules

- *lsmod | grep test*
- *sudo chmod +777 /dev/test_mod01*
- **sudo ./a.out**

Command	Description
gcc read_write_mod.c	compile user space application
make	compile kernel module
<i>lsmod grep "test"</i>	List modules, filter using string "test"
sudo insmod test_mod01.ko	Insert new compiled module to our kernel

Resources

- The Linux Kernel Module Programming Guide
 - <https://tldp.org/LDP/lkmpg/2.6/html/>
- Linux Device Drivers
 - <https://lwn.net/Kernel/LDD3/>
- Linux Kernel Sources
 - <https://github.com/torvalds/linux>
 - <https://www.kernel.org/>